

1. What is Artificial Intelligence (AI)?

AI is the simulation of human intelligence in machines to perform tasks like learning, reasoning, and problem-solving.

2. What are the types of AI?

Narrow AI, General AI, and Super AI based on capability.

3. What is machine learning (ML)?

A subset of AI where systems learn patterns from data without explicit programming.

4. What is supervised learning?

ML where models are trained on labeled data with known outputs.

5. What is unsupervised learning?

ML where models identify patterns from unlabeled data.

6. What is reinforcement learning?

ML where agents learn optimal actions through rewards and penalties.

7. What is a neural network?

A computational model inspired by the human brain to recognize patterns.

8. What are the main components of a neural network?

Input layer, hidden layer(s), output layer, and activation functions.

9. What is an artificial neuron?

A unit that processes inputs with weights, adds bias, and applies an activation function.

10. What is the activation function?

It determines the output of a neuron given its input; e.g., sigmoid, ReLU.

11. What is sigmoid function?

A smooth curve mapping input to a range of 0 to 1.

12. What is ReLU function?

Outputs zero for negative input and linear for positive input.

13. What is softmax function?

Converts output vector into probabilities summing to 1.

14. What is the difference between AI and ML?

AI is the broader concept, ML is a technique to achieve AI.

15. What is overfitting in ML?

When a model performs well on training data but poorly on unseen data.

16. What is underfitting in ML?

When a model fails to capture patterns, performing poorly on training and test data.

17. What is backpropagation?

Algorithm to update neural network weights using gradient descent.

18. What is gradient descent?

An optimization method to minimize the loss function by adjusting weights.

19. What is learning rate?

A parameter controlling the step size during weight updates.

20. What is a loss function?

A measure of how far predictions are from actual values.

21. What is mean squared error (MSE)?

Average squared difference between predicted and actual values.

22. What is cross-entropy loss?

A loss function measuring the difference between probability distributions.

23. What is a feedforward neural network?

Neural network where connections move forward from input to output.

24. What is a recurrent neural network (RNN)?

Neural network with cycles allowing sequential data processing.

25. What is a convolutional neural network (CNN)?

Network using convolutional layers for image and spatial data processing.

26. What is pooling in CNN?

Reduces spatial dimensions to decrease computation and control overfitting.

27. What is max pooling?

Selects the maximum value in a region of feature maps.

28. What is dropout in neural networks?

Randomly disables neurons during training to prevent overfitting.

29. What is batch normalization?

Normalizes layer inputs to stabilize learning and improve speed.

30. What is the difference between batch and stochastic gradient descent?

Batch uses all data, stochastic uses one sample per update.

31. What is a perceptron?

A simple linear binary classifier in neural networks.

32. What is multi-layer perceptron (MLP)?

A feedforward network with one or more hidden layers.

33. What is weight initialization?

Setting initial neuron weights to improve convergence.

34. What is vanishing gradient problem?

Gradients become too small, slowing or stopping learning in deep networks.

35. What is exploding gradient problem?

Gradients become too large, causing unstable updates.

36. What is a radial basis function (RBF) network?

Network using radial basis functions as activation functions for hidden neurons.

37. What is autoencoder?

A neural network that compresses input into a latent space and reconstructs it.

38. What is a generative adversarial network (GAN)?

A network with generator and discriminator competing to create realistic data.

39. What is reinforcement signal in AI?

A feedback signal representing reward or penalty for an agent's action.

40. What is Q-learning?

A model-free RL algorithm that learns optimal action-value function.

41. What is policy in reinforcement learning?

A strategy mapping states to actions for an agent.

42. What is value function in RL?

Expected cumulative reward for being in a state under a policy.

43. What is epsilon-greedy strategy?

Selects the best-known action most of the time, random action sometimes.

44. What is supervised vs reinforcement learning?

Supervised uses labeled data; RL learns from interaction and rewards.

45. What is feature extraction in AI?

Transforming raw data into useful inputs for a model.

46. What is dimensionality reduction?

Reducing features while preserving information, e.g., PCA.

47. What is PCA?

Principal Component Analysis reduces correlated features into uncorrelated components.

48. What is clustering in ML?

Grouping similar data points together without labels.

49. What is K-means clustering?

An iterative algorithm partitioning data into K clusters by distance.

50. _____

51. Advanced Level (50–98)

52. What is deep learning?

Subset of ML using deep neural networks for feature learning and prediction.

53. What is the difference between shallow and deep networks?

Shallow has one hidden layer, deep has multiple hidden layers.

54. What is transfer learning?

Using a pre-trained model for a new but similar task.

55. What is fine-tuning in transfer learning?

Adjusting the weights of a pre-trained model on new data.

56. What is attention mechanism in neural networks?

Allows the network to focus on important parts of input sequences.

57. What is transformer architecture?

Uses self-attention to process sequences in parallel efficiently.

58. What is the difference between RNN and LSTM?

LSTM has memory cells to avoid vanishing gradient in sequential data.

59. What is GRU?

Gated Recurrent Unit, a simplified LSTM with fewer parameters.

60. What is bidirectional RNN?

Processes sequence in both forward and backward directions.

61. What is word embedding?

Maps words to dense vector representations capturing semantics.

62. What is one-hot encoding?

Represents categorical data as binary vectors with a single 1.

63. What is NLP?

Natural Language Processing deals with human language understanding by machines.

64. What is sentiment analysis?

Determining emotions or opinions from text data.

65. What is sequence-to-sequence model?

Maps input sequences to output sequences, e.g., translation.

66. What is attention vs self-attention?

Attention focuses on input-output relations, self-attention focuses within input.

67. What is BERT?

Bidirectional Encoder Representations from Transformers for NLP tasks.

68. What is GPT?

Generative Pre-trained Transformer, generates human-like text.

69. What is reinforcement learning vs supervised learning?

RL learns via interaction and rewards; supervised learns from labeled data.

70. What is Markov Decision Process (MDP)?

Mathematical framework modeling RL problems with states, actions, rewards.

71. What is reward shaping in RL?

Modifying rewards to speed up learning without changing optimal policy.

72. What is exploration vs exploitation in RL?

Exploration tries new actions, exploitation chooses best-known actions.

73. What is neural style transfer?

Combines content of one image with style of another using CNNs.

74. What is dropout vs batch normalization?

Dropout prevents overfitting, batch normalization stabilizes learning.

75. What is gradient clipping?

Limits gradient values to avoid exploding gradients.

76. What is adversarial attack in AI?

Inputs designed to mislead neural networks.

77. What is explainable AI (XAI)?

Techniques making AI decisions interpretable to humans.

78. What is AI bias?

Systematic error due to skewed or unbalanced data.

79. What is ethical AI?

Developing AI respecting fairness, transparency, and accountability.

80. What is convolution in CNN?

Sliding filter over input to extract spatial features.

81. What is stride in convolution?

Number of pixels the filter moves per step.

82. What is padding in CNN?

Adding zeros around input to preserve dimensions.

83. What is receptive field in CNN?

Input region affecting a neuron's output.

84. What is vanishing gradient in deep networks?

Gradients become too small, halting learning in early layers.

85. What is residual network (ResNet)?

Uses skip connections to avoid vanishing gradient in deep networks.

86. What is GAN mode collapse?

Generator produces limited outputs failing to capture data diversity.

87. What is reinforcement learning exploration-exploitation dilemma?

Choosing between trying new actions or exploiting known rewards.

88. What is temporal difference learning?

RL method estimating value functions using bootstrapping.

89. What is policy gradient method?

Directly optimizes policy using gradient ascent on expected reward.

90. What is actor-critic model?

RL model with separate actor (policy) and critic (value) networks.

91. What is AI search problem?

Finding a sequence of actions from initial to goal state.

92. What is BFS and DFS in AI search?

Breadth-first explores level by level, depth-first explores deep branches first.

93. What is heuristic function?

Estimates cost from current state to goal in informed search.

94. What is A algorithm?*

Search algorithm using cost + heuristic to find optimal path.

95. What is constraint satisfaction problem (CSP)?

Problem where variables must satisfy specific constraints.

96. What is forward vs backward chaining?

Forward derives conclusions from facts, backward starts from goal.

97. What is knowledge representation in AI?

Encoding information about the world for reasoning.

98. What is expert system?

AI system using knowledge base and inference engine to solve problems.

99. What is fuzzy logic?

Handles reasoning with degrees of truth rather than binary.

100. What is reinforcement learning reward function?

Defines feedback guiding an agent to learn desired behavior.

101. What is a perceptron in AI?

A single-layer neural network used for binary classification.

102. What is multi-class classification?

Predicting one label from more than two possible categories.

103. What is a loss landscape in neural networks?

A visualization of loss function values across different weights.

104. What is stochastic gradient descent with momentum?

Enhances SGD by adding previous update velocity to current step.

105. What is RMSProp optimizer?

Adjusts learning rate adaptively based on recent gradient magnitudes.

106. What is Adam optimizer?

Combines momentum and RMSProp for faster convergence.

107. **What is overparameterization in deep learning?**

Having more parameters than training samples, increasing capacity.

108. **What is early stopping?**

Halting training when validation performance stops improving to prevent overfitting.

109. **What is L1 vs L2 regularization?**

L1 adds absolute weights, L2 adds squared weights to loss.

110. **What is weight decay?**

Another term for L2 regularization applied during training.

111. **What is feature scaling?**

Normalizing features to similar ranges for faster convergence.

112. **What is standardization?**

Subtract mean and divide by standard deviation for normalization.

113. **What is min-max normalization?**

Rescales features to a specific range, typically 0 to 1.

114. **What is one-hot vector representation?**

Binary vector with one “1” and rest zeros representing a class.

115. **What is embedding layer in neural networks?**

Maps categorical data to dense continuous vector space.

116. **What is tokenization in NLP?**

Splitting text into words, subwords, or characters for processing.

117. **What is sequence padding in NLP?**

Adding tokens to make sequences of equal length for batching.

118. **What is dropout rate?**

Fraction of neurons randomly disabled during training.

119. **What is Xavier/Glorot initialization?**

Weights initialized to maintain variance across layers.

120. **What is He initialization?**

Weights initialized for ReLU activations to avoid vanishing gradients.

121. **What is vanishing gradient in RNNs?**

Gradients shrink exponentially, making learning early time steps difficult.

122. **What is exploding gradient in RNNs?**

Gradients grow exponentially, causing unstable weight updates.

123. **What is gradient clipping?**

Restricting gradients to a maximum threshold to stabilize training.

124. **What is attention score?**

Weight assigned to input elements based on relevance to output.

125. **What is self-attention?**

Mechanism computing attention among all elements in a sequence.

126. **What is multi-head attention?**

Using multiple self-attention mechanisms in parallel for richer representation.

127. **What is positional encoding in transformers?**

Adds order information to input tokens since transformers lack recurrence.

128. **What is encoder-decoder architecture?**

Encoder processes input, decoder generates output sequence.

129. **What is teacher forcing in RNN training?**

Using true previous output as next input to speed up training.

130. **What is beam search?**

Heuristic search keeping top-k sequences at each decoding step.

131. **What is greedy decoding?**

Selecting the highest probability token at each step for output.

132. **What is reinforcement learning agent?**

Learner interacting with an environment to maximize reward.

133. **What is Markov property in RL?**

Future state depends only on current state and action.

134. **What is temporal difference error?**

Difference between predicted and actual reward for learning updates.

135. **What is Q-value in Q-learning?**

Expected cumulative reward for a state-action pair.

136. **What is eligibility trace?**

Mechanism in RL assigning credit to previous states/actions.

137. **What is policy iteration?**

Iteratively evaluating and improving policy until convergence.

138. **What is value iteration?**

Updating value function directly to find optimal policy.

139. **What is deep Q-network (DQN)?**

Uses neural networks to approximate Q-values in complex environments.

140. **What is target network in DQN?**

A separate frozen network to stabilize Q-value updates.

141. **What is experience replay?**

Storing and sampling past transitions to break correlation during training.

142. **What is actor-critic in RL?**

Actor selects actions; critic evaluates action quality.

143. **What is advantage function in RL?**

Measures how much better an action is compared to average policy.

144. **What is PPO algorithm?**

Proximal Policy Optimization, an RL method balancing exploration and stability.

145. **What is TRPO algorithm?**

Trust Region Policy Optimization, constrains policy updates for stability.

146. **What is Monte Carlo method in RL?**

Learning value functions from complete episodes of experience.

147. **What is bootstrapping in RL?**

Estimating value using current estimates instead of waiting for final outcome.

148. **What is experience replay buffer?**

Memory storing transitions for random sampling in RL training.

149. **What is deep deterministic policy gradient (DDPG)?**

RL algorithm for continuous action spaces using actor-critic.

150. **What is neural network architecture search (NAS)?**

Automated search for optimal network structure.

151. **What is model compression in deep learning?**

Reducing model size for faster inference on limited hardware.

152. What is pruning in neural networks?

Removing insignificant weights or neurons to reduce complexity.

153. What is quantization in neural networks?

Reducing precision of weights and activations for efficiency.

154. What is knowledge distillation?

Transferring knowledge from a large model to a smaller one.

155. What is batch size?

Number of samples processed in one forward/backward pass.

156. What is epoch in deep learning?

One complete pass through the entire training dataset.

157. What is validation set?

Used to tune model hyperparameters without affecting training.

158. What is test set?

Evaluates final model performance on unseen data.

159. What is confusion matrix?

Table showing true vs predicted labels in classification tasks.

160. What is precision?

Fraction of correctly predicted positive samples among all predicted positives.

161. What is recall?

Fraction of correctly predicted positive samples among all actual positives.

162. What is F1-score?

Harmonic mean of precision and recall.

163. What is ROC curve?

Plots true positive rate vs false positive rate at different thresholds.

164. What is AUC?

Area under ROC curve, measures classifier performance.

165. What is cross-validation?

Technique to evaluate model generalization by splitting data into folds.

166. What is k-fold cross-validation?

Dividing data into k subsets and rotating training/testing.

167. **What is stratified sampling?**

Ensures each class proportion is maintained in splits.

168. **What is bias-variance tradeoff?**

Balancing underfitting (bias) and overfitting (variance).

169. **What is hyperparameter tuning?**

Optimizing parameters like learning rate, batch size, number of layers.

170. **What is grid search?**

Exhaustive search over predefined hyperparameter combinations.

171. **What is random search?**

Randomly sampling hyperparameter combinations for optimization.

172. **What is Bayesian optimization?**

Uses probabilistic model to select promising hyperparameters.

173. **What is recurrent cell in RNN?**

Basic unit maintaining hidden state across time steps.

174. **What is exploding/vanishing gradient solution?**

Use gradient clipping, LSTM/GRU, or proper initialization.

175. **What is deep reinforcement learning?**

Combines deep neural networks with RL for complex tasks.

176. **What is curriculum learning in AI?**

Training model on simpler tasks first, then gradually harder ones.

177. **What is reward shaping?**

Providing intermediate rewards to guide learning in RL.

178. **What is imitation learning?**

Learning by observing expert demonstrations.

179. **What is curiosity-driven learning?**

RL where intrinsic motivation drives exploration.

180. **What is multi-agent RL?**

Multiple agents learning and interacting in a shared environment.

181. **What is hierarchical RL?**

Learning high-level policies to select sub-policies or actions.

182. **What is meta-learning in AI?**

“Learning to learn” by training models to adapt quickly to new tasks.

183. **What is few-shot learning?**

Learning tasks with very limited labeled examples.

184. **What is zero-shot learning?**

Performing tasks without any labeled examples using prior knowledge.

185. **What is continual learning in AI?**

Learning new tasks without forgetting previous ones.

186. **What is catastrophic forgetting?**

When a model forgets previous knowledge while learning new tasks.

187. **What is generative model?**

Models data distribution to generate new samples.

188. **What is discriminative model?**

Models boundary between classes to predict labels.

189. **What is Variational Autoencoder (VAE)?**

Generative model learning latent distribution for data generation.

190. **What is KL-divergence?**

Measures difference between two probability distributions.

191. **What is GAN discriminator?**

Classifier distinguishing real from generated samples.

192. **What is GAN generator?**

Creates samples trying to fool the discriminator.

193. **What is conditional GAN?**

GAN generating samples conditioned on input labels.

194. **What is cycle GAN?**

Transforms data between domains without paired examples.

195. **What is attention visualization?**

Shows which parts of input the model focuses on.

196. **What is explainable AI (XAI) technique?**

LIME or SHAP provide feature importance and model insights.

197. **What is AI ethics challenge?**

Ensuring fairness, transparency, and accountability in AI systems.

198. **What is bias-variance in neural networks?**

Balancing model complexity to avoid underfitting or overfitting.

199. **What is hyperparameter vs parameter in NN?**

Parameter learned during training; hyperparameter set before training.

200. **What is residual connection?**

Shortcut connecting input to output of a layer to ease learning.

201. **What is dense layer in NN?**

Fully connected layer connecting every input to every output neuron.

202. **What is global average pooling?**

Reduces spatial dimensions by averaging each feature map.

203. **What is LSTM forget gate?**

Controls which information from previous state to discard.

204. **What is LSTM input gate?**

Controls which new information to store in memory cell.

205. **What is LSTM output gate?**

Controls which memory to expose as output at each time step.

206. **What is exploding gradient solution in RNN?**

Use gradient clipping or normalized initialization.

207. **What is learning curve in AI?**

Plot of model performance vs training progress.

208. **What is reinforcement learning environment?**

Simulated system where an agent takes actions and receives rewards.

209. **What is policy network in RL?**

Neural network that outputs action probabilities.

210. **What is value network in RL?**

Predicts expected future reward of states or state-action pairs.

211. **What is Monte Carlo Tree Search (MCTS)?**

Search algorithm combining random sampling and tree exploration.

212. **What is AlphaGo architecture?**

Combines deep neural networks with MCTS for board game play.

213. **What is knowledge graph in AI?**

Graph representing entities and their relationships for reasoning.

214. **What is semantic similarity in NLP?**

Degree to which two texts have the same meaning.

215. **What is cosine similarity?**

Measures similarity between vectors using angle between them.

216. **What is BLEU score?**

Metric to evaluate machine-translated text against references.

217. **What is ROUGE score?**

Metric for evaluating text summarization quality.

218. **What is perplexity in language models?**

Measures how well a probability model predicts test data.

219. **What is gradient penalty in GANs?**

Regularization term to stabilize training.

220. **What is mode collapse in GANs?**

Generator produces limited outputs ignoring diversity.

221. **What is Wasserstein GAN?**

GAN variant using Wasserstein distance for stable training.

222. **What is contrastive learning?**

Learning representations by comparing positive and negative pairs.

223. **What is self-supervised learning?**

Learning without labels by creating surrogate tasks.

224. **What is masked language modeling?**

Predicting missing tokens in a sentence, used in BERT.

225. **What is next sentence prediction?**

Task predicting whether two sentences follow each other.

226. **What is sequence labeling?**

Assigning a label to each token in a sequence.

227. **What is entity recognition?**

Identifying named entities in text like persons, places, and organizations.

228. **What is word sense disambiguation?**

Determining correct meaning of a word based on context.

229. **What is beam width in beam search?**

Number of sequences kept at each decoding step.

230. **What is greedy vs beam search tradeoff?**

Greedy is fast but suboptimal; beam explores more for accuracy.

231. **What is transformer encoder?**

Processes input sequence and generates context-rich embeddings.

232. **What is transformer decoder?**

Generates output sequence using encoder context and self-attention.

233. **What is positional mask in transformers?**

Prevents attention to future tokens during decoding.

234. **What is masked multi-head attention?**

Attention that masks future tokens for autoregressive modeling.

235. **What is layer normalization?**

Normalizes inputs across features for each sample.

236. **What is residual dropout?**

Applies dropout to residual connection for regularization.

237. **What is knowledge distillation loss?**

Loss comparing teacher model outputs with student predictions.

238. **What is hierarchical attention?**

Attention applied at multiple levels, e.g., word and sentence.

239. **What is positional embedding vs positional encoding?**

Embedding is learned; encoding is fixed function of position.

240. **What is beam search diversity?**

Adding penalty to encourage varied outputs.

241. **What is catastrophic forgetting solution?**

Replay memory or regularization methods like EWC.

242. **What is attention heatmap?**

Visual showing attention weights over input tokens.

243. **What is sequence length truncation?**

Cutting sequences to a maximum length for batching.

244. **What is multi-label classification?**

Predicting multiple labels per instance simultaneously.

245. **What is label smoothing?**

Adjusting target labels to prevent overconfident predictions.

246. **What is teacher-student training?**

Student learns from soft targets provided by teacher network.

247. **What is curriculum RL?**

Training RL agents from easier tasks to complex tasks sequentially.

248. **What is policy entropy in RL?**

Encourages exploration by maximizing randomness of policy.

249. **What is advantage actor-critic (A2C)?**

Uses advantage function to reduce variance in policy gradient.

250. **What is residual connection in transformers?**

Adds input to output of attention/feedforward layer to ease gradient flow.

251. **What is pre-training in transformers?**

Learning general representations before fine-tuning on tasks.

252. **What is fine-tuning in transformers?**

Adjusting pre-trained model weights for a specific task.

253. **What is multi-task learning?**

Training a single model for multiple related tasks simultaneously.

254. **What is hyperparameter search space?**

Range of possible hyperparameter values to explore.

255. **What is zero-shot classification?**

Predicting classes unseen during training using knowledge transfer.

256. **What is few-shot classification?**

Learning from a very small number of labeled examples.

257. **What is continual learning in neural networks?**

Updating model on new tasks without forgetting old tasks.

258. **What is elastic weight consolidation (EWC)?**

Regularization to reduce forgetting in continual learning.

259. **What is adversarial training?**

Training on adversarial examples to improve robustness.

260. **What is model ensembling?**

Combining predictions of multiple models to improve accuracy.

261. **What is stacked autoencoder?**

Multiple autoencoders stacked to learn hierarchical features.

262. **What is denoising autoencoder?**

Trains to reconstruct input from corrupted versions for robust features.

263. **What is sparse autoencoder?**

Encourages sparse activation to learn meaningful features.

264. **What is neural architecture search (NAS)?**

Automated search for optimal neural network structures.

265. **What is attention visualization in transformers?**

Shows which tokens the model focuses on for predictions.

266. **What is token embedding vs segment embedding?**

Token embedding for word identity; segment embedding for sentence distinction.

267. **What is deep Q-learning vs vanilla Q-learning?**

Deep Q-learning uses neural networks for Q-value approximation.

268. **What is double DQN?**

Uses separate networks to reduce overestimation in Q-learning.

269. **What is prioritized experience replay?**

Samples transitions with high temporal-difference error more frequently.

270. **What is target network soft update?**

Gradually updating target network to stabilize learning.

271. **What is curriculum learning in NLP?**

Training models on easier sequences first before complex ones.

272. **What is masked token prediction?**

Predicting missing tokens in self-supervised NLP tasks.

273. **What is tokenization vs subword tokenization?**

Subword splits rare words into smaller units to handle unknown words.

274. **What is byte-pair encoding (BPE)?**

Algorithm for subword tokenization by merging frequent pairs.

275. **What is sentencepiece tokenizer?**

Language-independent subword tokenizer for NLP.

276. **What is transformer encoder-decoder attention?**

Decoder attends to encoder outputs for sequence generation.

277. **What is cross-attention?**

Attention between different sequences, e.g., decoder attends to encoder.

278. **What is self-supervised pre-training?**

Learning from unlabeled data using surrogate tasks.

279. **What is masked language model vs causal LM?**

Masked predicts missing tokens; causal predicts next token autoregressively.

280. **What is AI model calibration?**

Adjusting model outputs so predicted probabilities reflect true likelihood.

281. **What is temperature scaling?**

Softens model predictions to improve calibration.

282. **What is contrastive predictive coding?**

Learning representations by predicting future latent embeddings.

283. **What is siamese network?**

Network comparing similarity between two inputs.

284. **What is triplet loss?**

Loss encouraging anchor-positive pairs closer than anchor-negative.

285. **What is metric learning?**

Learning a distance function for similarity-based tasks.

286. **What is transformer block?**

Unit with multi-head attention, feedforward network, and residual connections.

287. **What is feedforward network in transformer?**

Applies linear transformations and activation to each position independently.

288. **What is multi-task fine-tuning?**

Fine-tuning model on multiple related tasks at once.

289. **What is unsupervised pre-training?**

Learning representations without labeled data.

290. **What is semi-supervised learning?**

Learning from small labeled dataset plus large unlabeled dataset.

291. **What is curriculum RL vs standard RL?**

Curriculum gradually increases task difficulty to improve learning.

292. **What is knowledge graph embedding?**

Mapping entities and relations into vector space for reasoning.

293. **What is Graph Neural Network (GNN)?**

Neural network operating on graph-structured data.

294. **What is message passing in GNN?**

Nodes exchange information with neighbors to update embeddings.

295. **What is graph attention network (GAT)?**

GNN using attention mechanism to weigh neighbors differently.

296. **What is node classification in GNN?**

Predicting label of nodes based on graph structure.

297. **What is link prediction in GNN?**

Predicting existence of edges between nodes.

298. **What is graph pooling?**

Aggregating node features to form graph-level representation.